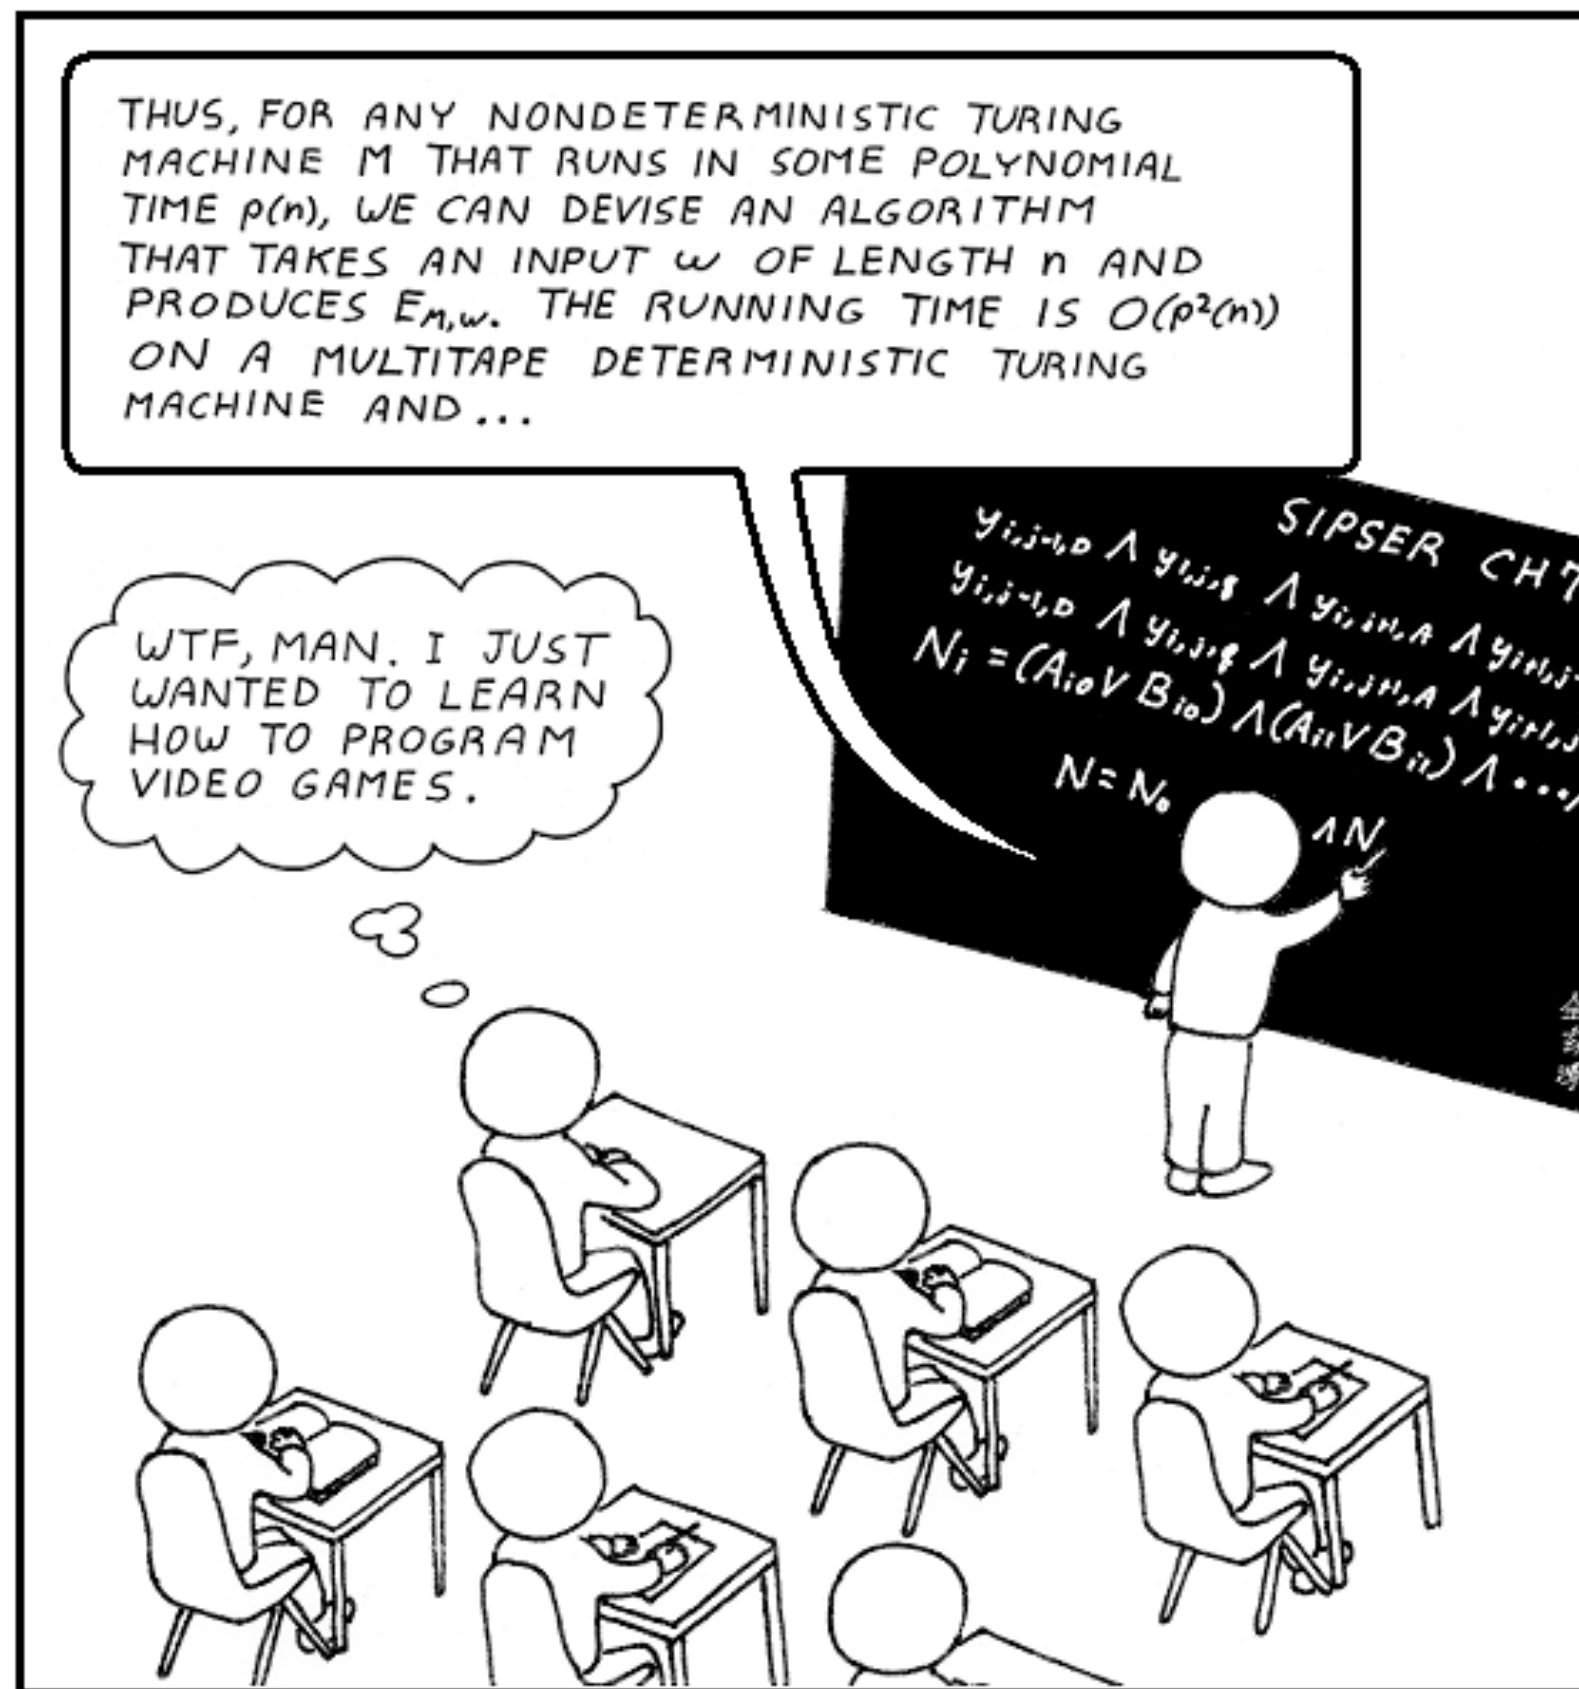


# CSE 549: Computational Biology

Computer Science for Biology

# What is Computer Science?



# What is Computer Science?

Not actually simple to define constructively

Still debate whether certain areas constitute CS

Computer science is the scientific and practical approach to computation and its applications. It is the systematic study of the feasibility, structure, expression, and mechanization of the methodical procedures (or algorithms) that underlie the acquisition, representation, processing, storage, communication of, and access to information\* ...

## What isn't Computer Science?

**Don't** install operating systems (may develop them)

**Don't** set up the office network (may study / design network protocols)

**Not about** Hacking together a program or learning a web-framework — programming  $\neq$  CS (may study formal languages and develop new programming languages, and programming is a skill many computer scientists learn / master.)

\*<http://www.cs.bu.edu/AboutCS/WhatIsCS.pdf>


# What is Computer Science?

Concerned with the development of provably **correct** and **efficient** computational procedures (algorithms & data structures) to answer **well-specified** problems.

To answer a computational question, we first need a well-formulated problem.


# Assembly

Reads



+

Reference genome



Input DNA



How to assemble  
puzzle without the  
benefit of knowing  
what the finished  
product looks like?

# Assembly

Whole-genome “shotgun” sequencing starts by copying and fragmenting the DNA

(“Shotgun” refers to the random fragmentation of the whole genome; like it was fired from a shotgun)

Input: GCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT


Copy: GCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT  
GCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT  
GCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT  
GCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT

Fragment: GCGTCTA TATCTCGG CTCTAGGCCCTC ATTTTTT  
GGC GTCTATAT CTCGGCTCTAGGCCCTCA TTTTTT  
GCGTC TATATCT CGGCTCTAGGCCCT CATTTTTT  
GCGTCTAT ATCTCGGCTCTAG GCCCTCA TTTTTT



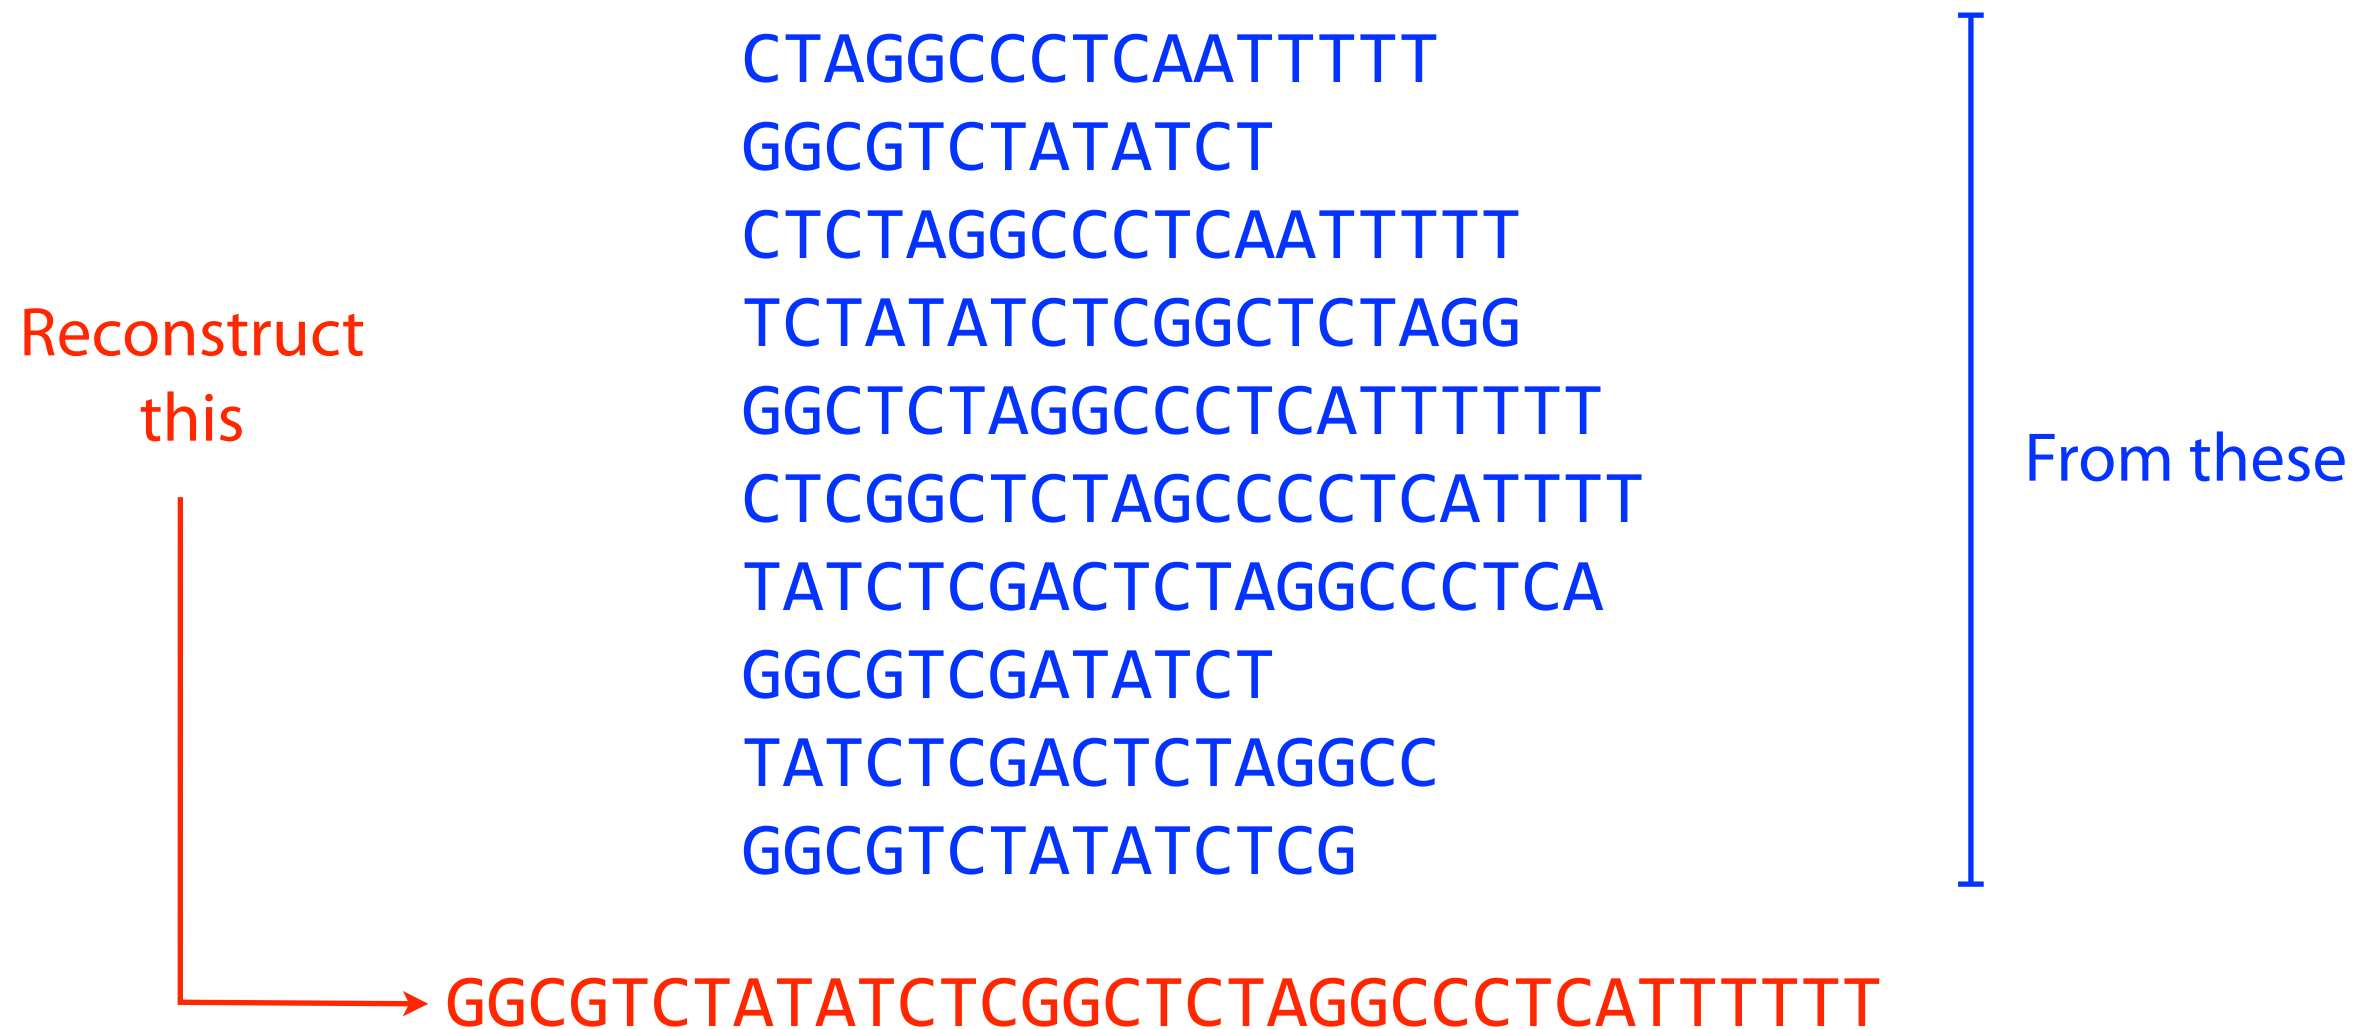
# Assembly

Assume sequencing produces such a large # fragments that almost all genome positions are *covered* by many fragments...



# Assembly

...but we don't know what came from where





# Assembly

CTAGGCCCTCAATTTT  
GGCGTCTATATCT  
CTCTAGGCCCTCAATTTT  
TCTATATCTCGGCTCTAGG  
GGCTCTAGGCCCTCATTTTT  
CTCGGCTCTAGCCCCTCATT  
TATCTCGACTCTAGGCCCTCA  
GGCGTCGATATCT  
TATCTCGACTCTAGGCC  
GGCGTCTATATCTCG



From  
these

Reconstruct this



??

# What is Computer Science?

Concerned with the development of provably **correct** and **efficient** computational procedures (algorithms & data structures) to answer **well-specified** problems.

To answer a computational question, we first need a well-formulated problem.

**Given:** a collection,  $R$ , of sequencing reads (strings)

**Find:** The genome (string),  $G$ , that generated them

# What is Computer Science?

Concerned with the development of provably **correct** and **efficient** computational procedures (algorithms & data structures) to answer **well-specified** problems.

To answer a computational question, we first need a well-formulated problem.

**Given:** a collection  $R$ , of sequencing reads (strings)

**Find:** The genome (string),  $G$ , that generated them

**Not** well-specified.

What makes one genome more likely than another?

What constraints do we place on the space of solutions?

# What is Computer Science?

Concerned with the development of provably **correct** and **efficient** computational procedures (algorithms & data structures) to answer **well-specified** problems.

To answer a computational question, we first need a well-formulated problem.

**Given:** a collection,  $R$ , of sequencing reads (strings)

**Find:** The shortest genome (string),  $G$ , that contains all of them



# Shortest Common Superstring

**Given:** a collection,  $S = \{s_1, s_2, \dots, s_k\}$  , of sequencing reads (strings)

**Find\*:** The shortest possible genome (string),  $G$ , such that  $s_1, s_2, \dots, s_k$  are all substrings of  $G$

**How**, might we go about solving this problem?

\*for reasons we'll explore later, this isn't actually a great formulation for genome assembly.

# Shortest common superstring

Given a collection of strings  $S$ , find  $SCS(S)$ : the shortest string that contains all strings in  $S$  as substrings

Without requirement of “shortest,” it’s easy: just concatenate them

Example:  $S$ : BAA AAB BBA ABA ABB BBB AAA BAB

Concatenation: BAAAABBBBAABAABBBBBBAAABAB  
└────────── 24 ─────────┘

$SCS(S)$ : AAABBBBABAA  
└── 10 ─┘

AAA  
AAB  
ABB  
BBB  
BBA  
BAB  
ABA  
BAA

Idea: pick order for strings in *S* and construct superstring

*order 1:* AAA AAB ABA ABB BAA BAB BBA BBB  
AAA



Idea: pick order for strings in *S* and construct superstring

*order 1:* AAA AAB ABA ABB BAA BAB BBA BBB  
AAAB

Idea: pick order for strings in  $S$  and construct superstring

*order 1:* AAA AAB ABA ABB BAA BAB BBA BBB  
          └──────────┘  
          AAABA

Idea: pick order for strings in *S* and construct superstring

*order 1:* AAA AAB ABA ABB BAA BAB BBA BBB  
AAABABB

Idea: pick order for strings in *S* and construct superstring

*order 1:* AAA AAB ABA ABB BAA BAB BBA BBB

AAABABBAABABBABBB ← superstring 1

Idea: pick order for strings in  $S$  and construct superstring

*order 1:* AAA AAB ABA ABB BAA BAB BBA BBB

AAABABBAABABBABBB ← superstring 1

*order 2:* AAA AAB ABA BAB ABB BBB BAA BBA

AAABABBBBAABBA ← superstring 2

Try all possible orderings and pick shortest superstring

If  $S$  contains  $n$  strings,  $n!$  ( $n$  factorial) orderings possible

*order 1:* AAA AAB ABA ABB BAA BAB BBA BBB

AAABABBAABABBABBB ← superstring 1

*order 2:* AAA AAB ABA BAB ABB BBB BAA BBA

AAABABBBBAABBA ← superstring 2

If  $S$  contains  $n$  strings,  $n!$  ( $n$  factorial) orderings possible

# Shortest common superstring

Can we solve it?

Imagine a modified overlap graph where each edge has cost = - (length of overlap)


SCS corresponds to a path that visits every node once, minimizing total cost along path

That's the *Traveling Salesman Problem (TSP)*, which is NP-hard!

S: AAA AAB ABB BBB BBA

SCS(S): AAABBBBA

AAA  
AAB  
ABB  
BBB  
BBA





# Shortest common superstring

Say we disregard edge weights and just look for a path that visits all the nodes exactly once


That's the *Hamiltonian Path* problem: NP-complete

Indeed, it's well established that SCS is NP-hard

$S$ : AAA AAB ABB BBB BBA

$SCS(S)$ : AAABBBBA

AAA  
AAB  
ABB  
BBB  
BBA



# Shortest common superstring & friends

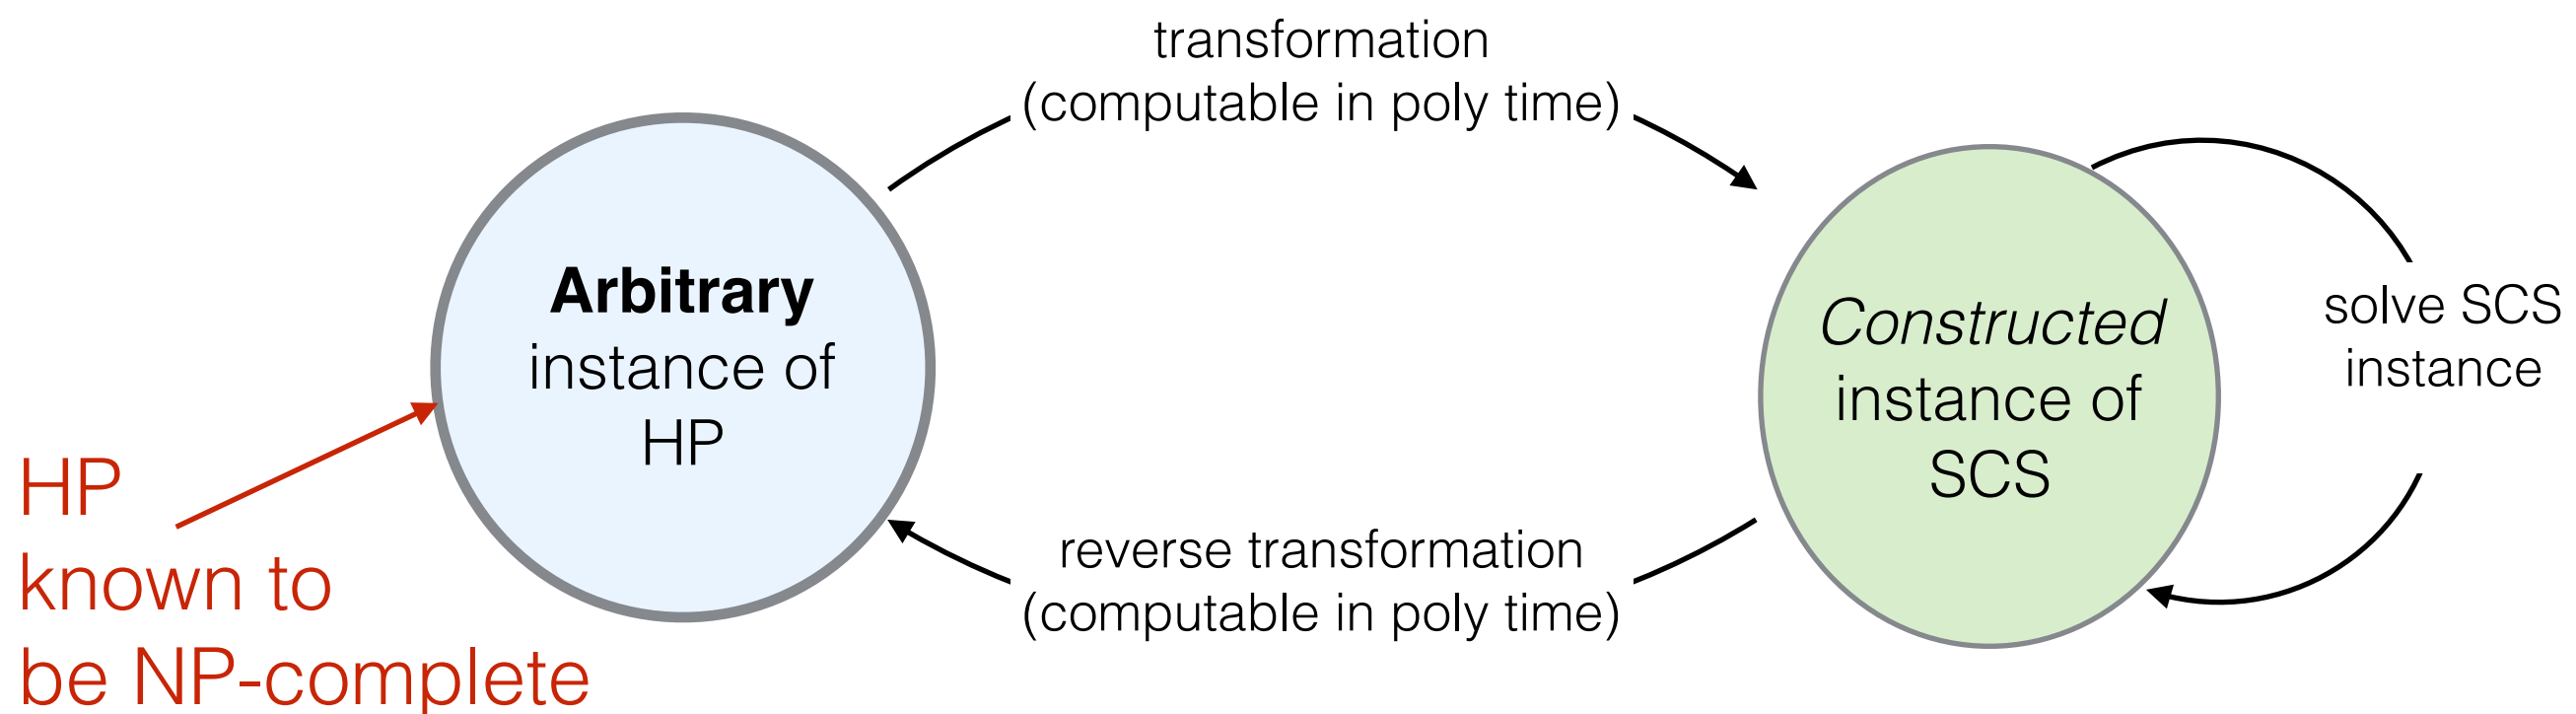
Traveling Salesman, Hamiltonian Path, and Shortest Common Superstring are all NP-hard

For refreshers on Traveling Salesman, Hamiltonian Path, NP-hardness and NP-completeness, see Chapters 34 and 35 of “Introduction to Algorithms” by Cormen, Leiserson, Rivest and Stein, or Chapters 8 and 9 of “Algorithms” by Dasgupta, Papadimitriou and Vazirani (free online: <http://www.cs.berkeley.edu/~vazirani/algorithms>)

**Important note:** The fact that we modeled SCS as NP-hard problems (TSP and HP) **does not** prove that (the decision version of) SCS is NP-complete. To do that, we must **reduce** a known NP-complete problem to **SCS**.

**Given** an instance  $I$  of a known hard problem, **generate** an instance  $I'$  of SCS such that if we can solve  $I'$  in polynomial time, then we can solve  $I$  in polynomial time. This *implies* that SCS is *at least* as hard as the hard problem.

This can be done e.g. with HAMILTONIAN PATH



# Shortest Common Superstring

The fact that SCS is **NP-complete** means that it is unlikely that there exists *any* algorithm that can solve a general instance of this problem in time polynomial in  $n$  — the number of strings.

If we give up on finding the *shortest* possible superstring  $G$ , how does the situation change?

# Shortest Common Superstring

There's a “greedy” *heuristic* that turns out to be an *approximation algorithm* (provides a solution within a constant factor of the optimum)


At *each step*, chose the *pair of strings* with the *maximum overlap*, merge them, and return the merged string to the collection.

*Greedy conjecture* factor of 2-OPT *is* the worst case


## Different approx. (**not all greedy**)

ratio	authors	year
approximating SCS		
3	Blum, Jiang, Li, Tromp and Yannakakis [4]	1991
$2\frac{8}{9}$	Teng, Yao [23]	1993
$2\frac{5}{6}$	Czumaj, Gasieniec, Piotrow, Rytter [8]	1994
$2\frac{50}{63}$	Kosaraju, Park, Stein [15]	1994
$2\frac{3}{4}$	Armen, Stein [1]	1994
$2\frac{50}{69}$	Armen, Stein [2]	1995
$2\frac{2}{3}$	Armen, Stein [3]	1996
$2\frac{25}{42}$	Breslauer, Jiang, Jiang [5]	1997
$2\frac{1}{2}$	Sweedyk [21]	1999
$2\frac{1}{2}$	Kaplan, Lewenstein, Shafrir, Sviridenko [12]	2005
$2\frac{1}{2}$	Paluch, Elbassioni, van Zuylen [18]	2012
$2\frac{11}{23}$	Mucha [16]	2013

# Greedy shortest common superstring




# Greedy shortest common superstring







# Greedy shortest common superstring




# Greedy shortest common superstring




# Greedy shortest common superstring




# Greedy shortest common superstring



# Greedy shortest common superstring



# Greedy shortest common superstring



# Greedy shortest common superstring

**AAABBBBA** ← superstring, length=7




# Shortest common superstring: greedy

Greedy-SCS: in each round, merge pair of strings with maximal overlap. Stop when there's 1 string left.  $l$  = minimum overlap.

Algorithm in action ( $l = 1$ ):

┌─── Input strings ──┐  
AAA AAB ABB BBB BBA



# Shortest common superstring: greedy


Greedy-SCS: in each round, merge pair of strings with maximal overlap. Stop when there's 1 string left.  $l$  = minimum overlap.

## Algorithm in action ( $l = 1$ ):

Input strings

AAA AAB ABB BBB BBA

AAA AAB ABB BBB BBA



# Shortest common superstring: greedy

Greedy-SCS: in each round, merge pair of strings with maximal overlap. Stop when there's 1 string left.  $l$  = minimum overlap.


Algorithm in action ( $l = 1$ ):

┌─── Input strings ──┐

AAA AAB ABB BBB BBA

AAA AAB ABB BBB BBA

AAAB ABB BBB BBA



# Shortest common superstring: greedy

Greedy-SCS: in each round, merge pair of strings with maximal overlap. Stop when there's 1 string left.  $l$  = minimum overlap.


Algorithm in action ( $l = 1$ ):

┌─── Input strings ──┐

AAA AAB ABB BBB BBA

AAA AAB ABB BBB BBA

AAAB ABB BBB BBA



# Shortest common superstring: greedy

Greedy-SCS: in each round, merge pair of strings with maximal overlap. Stop when there's 1 string left.  $l$  = minimum overlap.

Algorithm in action ( $l = 1$ ):


┌─── Input strings ──┐

AAA AAB ABB BBB BBA

AAA AAB ABB BBB BBA

AAAB ABB BBB BBA

AAAB BBBA ABB



# Shortest common superstring: greedy

Greedy-SCS: in each round, merge pair of strings with maximal overlap. Stop when there's 1 string left.  $l$  = minimum overlap.

Algorithm in action ( $l = 1$ ):


┌─── Input strings ──┐

AAA AAB ABB BBB BBA

AAA AAB ABB BBB BBA

AAAB ABB BBB BBA

AAAB BBBA ABB



# Shortest common superstring: greedy

Greedy-SCS: in each round, merge pair of strings with maximal overlap. Stop when there's 1 string left.  $l$  = minimum overlap.

Algorithm in action ( $l = 1$ ):

┌─── Input strings ──┐


AAA AAB ABB BBB BBA

AAA AAB ABB BBB BBA

AAAB ABB BBB BBA

AAAB BBBA ABB

AAABB BBBA



# Shortest common superstring: greedy

Greedy-SCS: in each round, merge pair of strings with maximal overlap. Stop when there's 1 string left.  $l$  = minimum overlap.

Algorithm in action ( $l = 1$ ):

┌─── Input strings ──┐

AAA AAB ABB BBB BBA

AAA AAB ABB BBB BBA

AAAB ABB BBB BBA

AAAB BBBA ABB

AAABB BBBA

AAABBBA

AAABBBA


That's the SCS




# Greedy shortest common superstring

AAA AAB ABB BBA BBB  
↓ ↙  
AAAB ABB BBA BBB


# Greedy shortest common superstring



# Greedy shortest common superstring

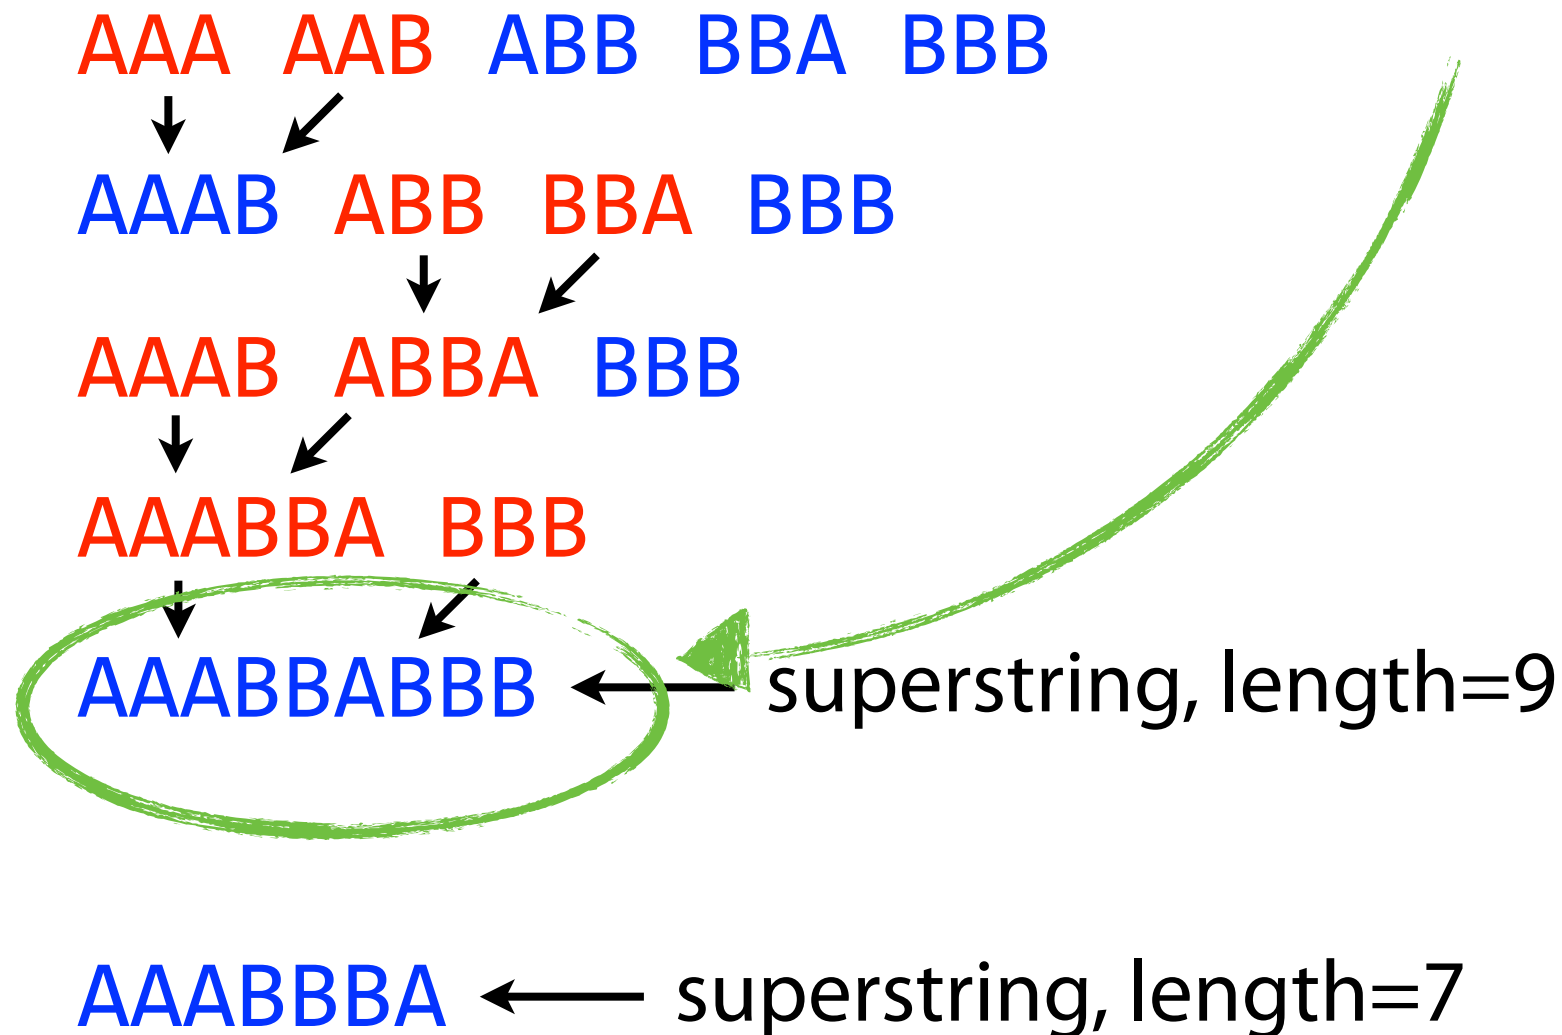


# Greedy shortest common superstring



# Greedy shortest common superstring

Note: approx. guarantee  
is on length of the superstring  
Actual result may be very different.



Greedy answer *isn't necessarily optimal*

# Take-home message:

We are interested in *correct and efficient algorithms* for solving *well-specified* problems.

We must be careful about how we *pose* the problems.

Actually, shortest common superstring is a rather poor model for sequence assembly, due to repeats and errors.